

Wardle Editorial Canon

Standards for Governed Editorial Artifacts

Part VI — AI Draft Standards

1. Purpose

AI Draft Standards define how artificial intelligence may contribute to the creation, analysis, revision, and maintenance of governed clinical knowledge within the Wardle Editorial Operating System.

Artificial intelligence may increase editorial capacity by:

- producing initial drafts
- proposing revisions
- identifying missing knowledge
- generating structured educational content
- suggesting relationships
- analysing reasoning quality
- detecting inconsistencies
- recommending editorial actions

Artificial intelligence does not possess editorial authority.

Its outputs begin as candidate knowledge and must pass through the same clinical, educational, and governance standards applied to manually authored work.

The purpose of AI governance is therefore not to restrict useful generation.

It is to ensure that AI contribution remains:

- clinically grounded
 - educationally purposeful
 - transparent
 - reviewable
 - reversible
 - attributable
 - subordinate to human judgement
-

2. Foundational Principle

AI may contribute to editorial production.

AI may not confer institutional trust.

The canonical flow is:

Governed Editorial Context

↓

AI Generation or Analysis

↓

Candidate Draft

↓

Automated Validation

↓

Human Editorial Review

↓

Accept, Reject, Request Changes, or Supersede

↓

Controlled Application

↓

Governed Artifact Revision

No stage before human editorial decision creates approved knowledge.

3. AI as an Editorial Contributor

Within WEOS, AI should be treated as a contributor with defined capabilities and limitations.

AI may:

- synthesise provided context
- structure content
- generate controlled variants
- detect likely omissions
- compare artifacts
- propose corrections
- identify possible contradictions
- generate candidate relationships
- assist with maintenance

AI must not be treated as:

- a clinical authority
- an autonomous reviewer
- a publisher
- an approver
- an evidence source
- a substitute for specialty judgement

The credibility of AI output depends on its governed context, validation, and human review.

4. AI Artifact Families

AI may generate or contribute to:

- Clinical Case Drafts
- Case Explanation Drafts
- Clue Revision Drafts
- Diagnosis Education Drafts
- Education Section Drafts
- Recall Prompt Drafts
- Teaching Rule Drafts
- Editorial Brief Drafts
- Graph Candidates
- Teaching Relationship Candidates
- Evidence Relationship Candidates
- Reasoning Path Drafts
- Differential Suggestions
- Unsupported-Claim Repairs
- Editorial Recommendations

Each AI output should be governed according to the standard of the artifact it proposes to create or revise.

There is no separate reduced clinical standard for AI-generated content.

5. AI Draft Definition

An AI Draft is a proposed artifact or proposed change produced with material artificial-intelligence involvement.

An AI Draft may be:

- a complete new artifact
- a replacement section
- a focused revision
- an added clue
- a relationship proposal
- a correction to an unsupported claim
- a structured recommendation
- an alternative formulation

An AI Draft is not the governed artifact itself.

It remains separate until accepted and applied through an explicit editorial action.

6. AI Provenance

6.1 Purpose

AI provenance records how and why a draft was produced.

It allows reviewers to understand:

- what prompted generation
- what context was supplied
- what the system produced
- which artifact was affected
- what decision followed
- whether the output was applied

Without provenance, AI-assisted editorial change becomes difficult to audit or reproduce.

6.2 Required provenance

Every material AI Draft should preserve:

- diagnosis identity
 - generation purpose
 - action type
 - source issue or identified need
 - input context
 - source artifacts
 - source artifact versions
 - applied teaching rules
 - reasoning path where relevant
 - generated output
 - affected artifact type
 - affected artifact identity
 - generation time
 - model or generator version where available
 - validation outcome
 - reviewer decision
 - reviewer identity
 - review note
 - application status
 - resulting artifact revision where applied
-

6.3 Source issue

The source issue explains why AI generation occurred.

Examples include:

- missing Diagnosis Education
- weak investigation section
- incomplete differential coverage
- premature clue leakage
- unsupported management claim
- absent recall prompts
- insufficient case inventory
- missing escalation evidence
- inconsistent explanation
- outdated scoring-system guidance

AI generation should respond to an identifiable editorial need.

Unbounded generation without purpose should be avoided.

6.4 Input context

Input context is the governed knowledge supplied to the AI system.

It may contain:

- diagnosis metadata
- editorial brief
- learning goals
- teaching rules
- reasoning paths
- active graph facts
- evidence relationships
- approved cases
- approved Diagnosis Education
- required mimics
- discriminators
- pitfalls
- investigations
- scoring systems
- management anchors
- difficulty guidance
- reviewer instructions

The exact context used should remain reconstructable.

7. Generation Context

7.1 Purpose

Generation Context defines the governed editorial boundaries within which AI may draft content.

It should reduce the likelihood that the AI:

- invents educational priorities
 - ignores required mimics
 - contradicts teaching rules
 - introduces unsupported findings
 - reveals a diagnosis too early
 - creates generic education
 - recommends unsafe management
 - generates at the wrong difficulty
-

7.2 Context components

A generation context may include:

- diagnosis
 - concise clinical context
 - learning goals
 - mandatory content
 - prohibited or discouraged content
 - mimics
 - discriminators
 - pitfalls
 - investigations
 - scoring systems
 - management anchors
 - required teaching units
 - acceptable manifestations
 - editorial brief
 - difficulty strategy
 - source-availability summary
-

7.3 Diagnosis context

The diagnosis context should identify:

- canonical diagnosis
- display label
- specialty
- category
- body system
- clinical setting
- baseline difficulty
- age group where relevant
- urgency where relevant

AI should not infer diagnosis identity from an ungoverned descriptive title where a registry identity exists.

7.4 Learning goals

Learning goals define what the generated artifact should teach.

The AI should not introduce a competing educational objective unless explicitly permitted.

For example, a case intended to teach differential discrimination should not become primarily a management-recall case.

7.5 Mandatory content

Mandatory content may include:

- required findings
- required teaching points
- required mimics
- required investigations
- required management principles
- required complication recognition
- required escalation evidence

Mandatory content should be relevant to the artifact type.

Not every item must appear verbatim.

Its educational meaning must be preserved.

7.6 Avoidance guidance

Avoidance guidance identifies content or patterns the AI should not use.

Examples include:

- forbidden early clues
- pathognomonic wording in the opening
- unsupported diagnostic certainty
- generic educational phrases
- inappropriate investigation dependence
- unsafe management generalisation
- repetitive manifestations
- prohibited evidence patterns
- outdated terminology

Avoidance guidance should include reasons where possible.

7.7 Mimics and discriminators

Where diagnostic reasoning is required, generation context should identify:

- important mimics
- why they are confused
- key separators
- findings that favour the target
- findings that weaken the target
- mimics that should remain plausible until a later stage

AI should not manufacture artificial differentials merely to populate a section.

7.8 Difficulty strategy

Difficulty guidance may define:

- target difficulty
- expected diagnostic lock-in point
- findings to avoid too early
- differentials to keep alive
- acceptable alternative manifestations
- depth of investigation
- reasoning complexity

Difficulty should arise from controlled uncertainty and discrimination.

It should not arise from obscure terminology or withheld essential information.

7.9 Context sufficiency

Generation should not proceed automatically where required context is critically incomplete.

Possible blockers include:

- unresolved diagnosis identity
- diagnosis prohibited from generation
- absent educational scope
- contradictory teaching rules
- unresolved required mimic
- insufficient reasoning foundation
- no reliable management anchor for a management-focused artifact

Where generation proceeds with incomplete non-critical context, the limitation should be disclosed to reviewers.

8. Generation Purpose

Every AI generation request must identify its purpose.

Canonical purposes include:

- Case Generation
- Diagnosis Education Generation
- Education Section Regeneration
- Teaching Rule Generation
- Editorial Brief Generation
- Graph Candidate Generation
- Teaching Relationship Generation
- Evidence Relationship Generation
- Recall Prompt Generation
- Clue Revision
- Unsupported-Claim Repair
- Maintenance Update
- Quality Recommendation

The purpose determines:

- required context
 - output schema
 - validation standard
 - review packet
 - permitted application action
-

9. Structured Output

9.1 Purpose

AI output should use a defined structure appropriate to the target artifact.

Structured output improves:

- validation
- review
- traceability
- completeness assessment

- controlled application
 - revision comparison
-

9.2 Schema compliance

An AI Draft should satisfy the expected artifact schema before entering substantive review.

Schema failure may include:

- missing required fields
- invalid types
- invalid lifecycle value
- malformed identifiers
- duplicate stable keys
- incomplete structured entries
- content outside the permitted artifact shape

Schema compliance does not establish clinical quality.

9.3 No hidden narrative

Important claims should not be hidden within unstructured text where the artifact schema expects explicit fields.

For example:

- a discriminator should appear as a discriminator
 - a management caution should appear as a caution
 - an escalation implication should be explicitly represented
 - a scoring-system limitation should not be buried in a general paragraph
-

10. AI-Generated Clinical Case Draft

10.1 Purpose

An AI-Generated Clinical Case Draft proposes a new reasoning experience for a governed diagnosis.

It must comply with the Clinical Case Standards in Part III.

10.2 Required generation context

Case generation should consider:

- diagnosis identity
 - case purpose
 - learning goal
 - reasoning objective
 - target difficulty
 - required mimics
 - discriminators
 - teaching rules
 - forbidden evidence patterns
 - preferred clue types
 - excluded clue types
 - expected clue distribution
 - escalation requirements
 - management anchors where relevant
-

10.3 Case draft requirements

The draft should include:

- clinically plausible patient context
 - coherent chronology
 - appropriate opening uncertainty
 - structured clue sequence
 - meaningful differentials
 - defensible final diagnosis
 - aligned explanation
 - intended difficulty
 - source and generation provenance
-

10.4 Case-specific AI risks

Review should specifically assess for:

- premature leakage
- internally inconsistent findings
- invented laboratory values
- unrealistic chronology
- descriptive rather than canonical diagnosis
- incorrect clue explanations
- artificial differentials

- unsupported management
 - excessive narrative
 - clue duplication
 - failure to meet the intended learning goal
-

10.5 Case draft status

An AI-generated case should begin as a draft.

It must not enter approved, playable, ready-to-publish, or published states directly from generation.

11. AI-Generated Case Explanation

An AI-generated explanation should:

- interpret the actual clues
- identify how probability changed
- explain the major discriminator
- compare relevant mimics
- avoid invented findings
- preserve appropriate uncertainty
- use the canonical diagnosis
- align with the case revision reviewed

An explanation generated from an outdated case revision must not be applied to a newer revision without renewed assessment.

12. AI-Generated Clue Revision

12.1 Purpose

An AI-Generated Clue Revision proposes a focused repair to an existing case.

It may:

- rewrite a clue
- remove a clue
- add a clue
- reposition a clue
- narrow an over-specific clue
- strengthen a discriminator

- correct clinical inconsistency
-

12.2 Required content

A Clue Revision Draft should preserve:

- original clue
 - proposed clue
 - clue position
 - source issue
 - rationale
 - expected effect
 - case revision
 - validation findings
 - review status
-

12.3 Expected effect

The expected effect should be testable.

Examples include:

- reduce premature leakage
 - maintain pulmonary tuberculosis as a plausible mimic until clue four
 - correct the mismatch between the fever clue and its explanation
 - add evidence of neurological localisation
 - remove an unsupported laboratory abnormality
 - improve escalation recognition
-

12.4 Application

Accepting a Clue Revision Draft should not silently change the governed case.

Application should:

- identify the accepted draft
 - create or update an appropriate case revision
 - preserve the original clue
 - record the responsible editor
 - trigger revalidation where required
 - require renewed review where the reasoning changed materially
-

13. AI-Generated Diagnosis Education

13.1 Purpose

An AI-Generated Diagnosis Education Draft proposes a structured learner-facing resource for one diagnosis.

It must comply with the Diagnosis Education Standards in Part IV.

13.2 Required context

Education generation should consider:

- diagnosis identity
 - editorial brief
 - learning goals
 - teaching rules
 - clinical pattern
 - required findings
 - required mimics
 - discriminators
 - investigations
 - scoring systems
 - management anchors
 - complications
 - required pitfalls
 - recall concepts
 - references or evidence requirements
-

13.3 Section completeness

A complete generated education draft should address required sections through the expected structured schema.

Missing non-applicable content should be represented deliberately rather than concealed through generic filler.

13.4 Diagnosis specificity

AI-generated education must be checked for generic content that could apply unchanged to multiple unrelated diagnoses.

Weak content includes statements such as:

- may present with several symptoms
- diagnosis requires clinical correlation
- management depends on severity
- early diagnosis is important
- complications can be serious

These statements require diagnosis-specific clinical meaning.

13.5 Educational reasoning

Generated entries should explain:

- why a finding matters
- how it changes suspicion
- what it distinguishes
- what action it prompts
- which trap it prevents

Listing correct facts is insufficient.

14. AI-Generated Education Section

14.1 Purpose

Section generation allows a specific Diagnosis Education section to be drafted or repaired without regenerating the entire artifact.

Appropriate targets include:

- Summary
- Clinical Pattern
- Key Symptoms
- Key Signs
- Examination Pearls
- Scoring Systems
- Investigations
- Differentials
- Management
- Complications
- Pitfalls
- Recall Prompts
- References

14.2 Section context

Section generation should receive enough surrounding context to preserve coherence.

For example:

- Management generation should consider complications, urgency, and pitfalls.
- Differential generation should consider the clinical pattern and discriminators.
- Recall generation should use approved source sections.
- Scoring-system generation should consider investigation and management use.
- Pitfall generation should consider known mimics and teaching rules.

14.3 Regeneration safeguards

Section regeneration must preserve:

- source education version
- original section
- generated replacement
- reason for regeneration
- cross-section dependencies
- reviewer decision
- applied education revision

A regenerated section should not replace published content without controlled review and application.

15. AI-Generated Recall Prompt

An AI-generated Recall Prompt should:

- derive from governed education
- identify its source section
- test one meaningful concept
- provide a concise answer
- include an explanatory layer
- match the declared difficulty
- avoid unsupported extension beyond source content

AI should not produce recall prompts from unapproved or contradictory content without clearly marking the source as candidate knowledge.

16. AI-Generated Teaching Rule

16.1 Purpose

An AI-Generated Teaching Rule proposes a reusable diagnosis-specific instructional constraint.

It should not be created merely by converting arbitrary prose into rule form.

16.2 Required structure

The draft should propose:

- stable-key candidate
 - title
 - category
 - importance
 - rationale
 - acceptable manifestations
 - required differentials
 - expected evidence
 - difficulty guidance
 - timing constraint
 - applicability
-

16.3 Review focus

Reviewers should determine:

- whether the rule is truly reusable
- whether it reflects clinical and educational intent
- whether it duplicates an existing rule
- whether its importance is justified
- whether it should block generation or publication
- whether it applies to cases, education, graph work, or more than one domain

AI-generated teaching rules must not become active automatically.

17. AI-Generated Editorial Brief

An AI may propose a Diagnosis Editorial Brief where no brief exists or where significant repair is required.

The draft may include:

- editorial summary
- learning goals
- required teaching rules
- required mimics
- required pitfalls
- key investigations
- management anchors
- difficulty guidance
- case-generation guidance
- education guidance
- graph guidance

Because the Editorial Brief defines downstream educational intent, its review should occur before it is used as authoritative generation context.

18. AI-Generated Graph Candidate

18.1 Purpose

AI may extract or propose structured knowledge from cases and education.

All generated graph knowledge must begin as a Graph Candidate.

18.2 Required provenance

The candidate should preserve:

- source artifact
 - source version
 - source path
 - raw text
 - normalised text
 - candidate type
 - target diagnosis where relevant
 - confidence
 - structured payload
-

18.3 Prohibited automatic promotion

AI must not automatically convert an extracted statement into an active Graph Fact.

The candidate must undergo:

- identity resolution
 - duplicate assessment
 - scope assessment
 - clinical review
 - provenance confirmation
 - promotion decision
-

19. AI-Generated Teaching Relationship

AI may propose a relationship between two diagnoses.

The proposal should identify:

- source diagnosis
- target diagnosis
- relationship type
- teaching purpose
- shared presentation
- discriminator
- confusion reason
- learner pitfall
- supporting artifacts
- confidence

Reviewers must assess directionality.

A relationship that is useful from Diagnosis A to Diagnosis B may require different content in the reverse direction.

20. AI-Generated Evidence Relationship

AI may propose how an Evidence Node relates to a diagnosis.

The proposal should identify:

- evidence node
- diagnosis
- relationship type
- strength
- discriminator weight
- reasoning summary

- contradictory diagnoses
- source support

AI should not overuse absolute relationships such as `RULES_OUT`.

All proposed relationships begin as candidates.

21. AI-Generated Reasoning Path

An AI-generated Reasoning Path may propose:

- reasoning goal
- primary differentials
- supporting relationships
- discriminator evidence
- escalation evidence
- contradictory evidence
- required teaching points
- forbidden evidence patterns
- clue distribution
- generation purpose

A generated Reasoning Path must not guide production generation until reviewed and activated.

22. AI-Generated Differential Suggestion

AI may suggest differentials for:

- a clinical case
- Diagnosis Education
- an Editorial Brief
- a Reasoning Path
- a Teaching Relationship

Each suggestion should be evaluated for:

- clinical plausibility
- diagnosis identity
- educational relevance
- role
- discriminator availability
- safety significance
- duplicate risk

A suggested diagnosis should not be linked automatically where identity is ambiguous.

23. Unsupported-Claim Repair

23.1 Purpose

AI may propose a repair where a claim has been flagged as unsupported, unsafe, contradictory, or overstated.

23.2 Repair strategies

A repair may:

- remove the claim
 - narrow its scope
 - qualify certainty
 - add missing context
 - align it with governed evidence
 - replace unsafe management guidance
 - add a supporting reference candidate
 - separate universal from subgroup-specific claims
-

23.3 Review requirement

The repair must be reviewed against:

- original claim
- source issue
- evidence context
- dependent sections
- intended learner meaning
- potential clinical consequence

AI-generated repair does not prove the revised claim is correct.

24. AI Draft Validation

24.1 Purpose

AI Draft Validation evaluates whether a generated draft satisfies defined structural, clinical-reasoning, educational, and governance requirements before or during human review.

Validation should reduce reviewer burden.

It should not replace reviewer judgement.

24.2 Validation dimensions

Validation may assess:

- schema compliance
 - diagnosis alignment
 - teaching-rule compliance
 - reasoning coverage
 - evidence coverage
 - discriminator coverage
 - differential quality
 - hallucination risk
 - unsupported claims
 - contradiction risk
 - genericity
 - management safety
 - difficulty alignment
 - clue progression
 - cross-section consistency
 - publication blockers
-

24.3 Validation outputs

Validation may produce:

- validation status
- trust score
- trust tier
- blockers
- warnings
- strengths
- hallucination-risk signals

- reasoning-coverage analysis
 - evidence-coverage analysis
 - discriminator-coverage analysis
 - unsupported-claim signals
 - recommendations
 - validator version
-

24.4 Trust score

A Trust Score is a structured assessment of draft reliability relative to known context and rules.

It is not a probability that the content is medically correct.

It must not be presented as equivalent to clinical approval.

24.5 Trust tiers

Canonical Trust Tiers include:

High Trust

The draft demonstrates strong alignment with known context and no detected critical blocker.

Human review remains required.

Review Required

The draft is potentially usable but contains issues or uncertainty requiring focused human judgement.

Low Trust

The draft contains substantial weakness, missing support, or reasoning inconsistency.

It should not be applied without major revision.

Blocked

The draft contains a critical failure preventing safe progression.

Examples include:

- unresolved diagnosis identity
- unsafe management guidance

- major contradiction
 - insufficient required context
 - severe hallucination risk
 - failed schema
 - prohibited content pattern
-

24.6 Validation status

Canonical statuses include:

- Passed
- Needs Review
- Failed

A passed status means no configured blocking failure was detected.

It does not mean the draft is approved.

25. Hallucination-Risk Signals

25.1 Purpose

Hallucination-risk signals identify content that may have been invented, misapplied, overstated, or disconnected from the supplied context.

25.2 Common signals

Signals may include:

- finding absent from source context
 - invented scoring system
 - fabricated threshold
 - unsupported drug dose
 - nonexistent relationship
 - false diagnostic exclusivity
 - invented reference
 - incorrect guideline attribution
 - explanation of a clue that does not exist
 - management recommendation outside the diagnosis scope
 - contradictory disease mechanism
-

25.3 Signal interpretation

A hallucination-risk signal is not always proof of error.

It should guide human verification.

High-risk claims with potential learner harm should block application until resolved.

26. AI Draft Review Packet

Every AI Draft review should provide the reviewer with a complete packet containing:

- draft purpose
- source issue
- target artifact
- diagnosis context
- original content where revising
- generated content
- generation context
- applied rules
- reasoning path where relevant
- validation outcome
- blockers
- warnings
- unsupported-claim signals
- differences from current content
- expected effect
- available decisions
- downstream impact

Reviewers should not be required to infer why the draft exists.

27. AI Draft Review

27.1 Purpose

AI Draft Review is the human editorial evaluation of an AI-generated proposal.

It determines whether the proposal should be:

- accepted
- rejected

- revised
 - superseded
 - applied
 - deferred
-

27.2 Review dimensions

The reviewer should assess:

- clinical accuracy
 - diagnosis alignment
 - educational value
 - artifact-standard compliance
 - reasoning quality
 - source adequacy
 - evidence sufficiency
 - specificity
 - internal consistency
 - learner safety
 - downstream effects
 - AI validation findings
-

27.3 Review independence

Reviewers should not assume that an AI draft is reliable because:

- it is well written
- it follows the schema
- validation passed
- confidence is high
- the model is advanced
- similar content exists elsewhere

Fluency must not be confused with authority.

28. AI Draft Lifecycle

The canonical AI Draft lifecycle is:

Draft

↓

Review Required

↓

Pending Review

↓

Accepted, Rejected, or Needs Changes

↓

Applied or Superseded

28.1 Draft

The AI output has been created but may not yet satisfy minimum review requirements.

28.2 Review Required

The draft has been identified as needing human editorial evaluation.

28.3 Pending Review

The draft is queued or prepared for a reviewer.

28.4 Accepted

The reviewer has determined that the proposal is suitable for controlled application.

Acceptance does not itself mean that application has occurred.

28.5 Rejected

The proposal has not been accepted.

The draft and rationale should remain traceable.

28.6 Needs Changes

The proposal may become acceptable after revision.

The requested changes should be actionable.

28.7 Superseded

The proposal has been replaced by another draft, direct editorial change, or newer artifact version.

Supersession should identify the replacement where possible.

29. Acceptance and Application

Acceptance and application are distinct.

Acceptance

The reviewer approves the proposed content or change.

Application

The accepted proposal is incorporated into the governed artifact through a controlled revision.

The separation is necessary because:

- the underlying artifact may have changed since generation
 - dependencies may have become outdated
 - multiple accepted drafts may conflict
 - application may require a new revision
 - revalidation may be required
-

30. Application Preconditions

Before applying an accepted AI Draft, WEOS should verify:

- target artifact still exists
- target version matches the reviewed version
- diagnosis identity remains unchanged
- no conflicting revision has occurred
- required permissions are present
- critical blockers remain resolved
- dependent artifacts are known
- accepted output has not been superseded
- application will produce an auditable revision

Where the target has changed materially, the draft should return for comparison or regeneration.

31. Controlled Application

Application should preserve:

- accepted draft identity
- reviewer
- application actor
- target artifact
- previous version
- resulting version
- applied content
- application time
- validation triggered
- downstream review requirements

Application must not overwrite historical content invisibly.

32. Revalidation After Application

An applied AI Draft may require revalidation where it changes:

- clinical claims
- diagnosis mapping
- case clues
- clue order
- differential relationships
- explanation
- management
- scoring systems
- reasoning path
- learner-facing content

Acceptance of the draft does not waive validation of the resulting artifact.

33. Downstream Impact

AI-assisted changes may affect related artifacts.

Examples include:

- changing a differential may affect Teaching Relationships
- changing a clue may affect Clue Progression Analysis
- changing Management may affect Pitfalls and Recall Prompts
- changing a scoring system may affect Investigations
- changing a Teaching Rule may affect existing cases
- changing a Graph Fact may affect generation context
- changing the Editorial Brief may affect diagnosis readiness

The application process should surface known downstream effects.

34. AI Draft Conflict

A conflict occurs when:

- two drafts propose incompatible changes
- the target artifact has changed
- a newer draft supersedes an older one
- the draft contradicts an approved Teaching Rule
- an accepted draft conflicts with updated evidence
- applying one draft invalidates another

Conflicts should be resolved explicitly.

The system should not apply drafts in creation order without evaluating compatibility.

35. AI Draft Supersession

An AI Draft should be superseded where:

- a newer draft replaces it
- the underlying issue has been resolved manually
- the target artifact has materially changed
- the generation context is outdated
- evidence or guidance has changed
- the draft no longer reflects current editorial intent

Superseded drafts should remain historically traceable but unavailable for ordinary application.

36. AI Feedback and Learning

Reviewer decisions may inform future AI-assistance quality.

Useful feedback includes:

- accepted without change
- accepted with modification
- rejected for clinical inaccuracy
- rejected for genericity
- rejected for duplication
- rejected for unsafe management
- rejected for reasoning mismatch
- rejected for poor difficulty
- superseded by newer context

This feedback may improve prompts, validators, context construction, and recommendation policies.

Reviewer decisions should not be used to create hidden autonomous policy changes without governance.

37. Human Authority

Human editors retain authority over:

- clinical correctness
- educational intent
- artifact approval
- relationship activation
- diagnosis mapping in ambiguous cases
- acceptance of evidence
- management safety
- publication
- retirement
- institutional trust

AI may recommend.

Human reviewers decide.

38. Role and Permission Governance

AI Draft actions should be restricted according to editorial role.

Examples include:

- contributors may request generation
- editors may review ordinary drafts
- specialty reviewers may approve high-risk clinical content
- senior editors may approve publication-affecting repairs
- administrators may apply controlled migrations
- only authorised roles may publish resulting artifacts

Permission should reflect the consequence of the decision rather than the apparent simplicity of the draft.

39. High-Risk AI Content

Some AI outputs require heightened review.

High-risk content includes:

- drug doses
- emergency management
- procedural guidance
- contraindications
- paediatric treatment
- pregnancy-related management
- critical-care thresholds
- scoring-system cut-offs
- antimicrobial recommendations
- claims that a finding rules out a diagnosis
- guideline-dependent recommendations

High-risk content may require:

- explicit reference support
 - specialty review
 - dual review
 - stronger validation
 - publication blocking until resolved
-

40. AI and Evidence

AI is not an evidence source.

An AI-generated claim may be supported by:

- governed evidence relationships
- approved references
- recognised clinical guidelines
- verified primary or secondary evidence

The statement “generated by AI” provides provenance but no clinical authority.

AI should not fabricate, approximate, or imply references that were not verified.

41. AI and Publication

No AI-generated artifact should become learner-facing solely because:

- it passed schema validation
- it received a high trust score
- it was generated from approved context
- it resembles existing content
- no automated blocker was detected

Publication requires:

- artifact-level human approval
 - resolved blockers
 - appropriate evidence review
 - controlled application
 - review of the resulting artifact version
 - explicit publication decision
-

42. AI Operational Failure

AI generation may fail because of:

- provider unavailability
- timeout
- malformed output
- schema failure
- insufficient context
- safety restriction
- conflicting generation request
- output truncation
- validation failure

Operational failure should not alter the target artifact.

The system should preserve:

- attempted action
- failure type
- time
- diagnosis context
- retry eligibility
- absence of application

A failed generation must not leave an artifact in an ambiguous partially updated state.

43. AI Draft Maintenance

AI Drafts should be reviewed for staleness where:

- the target artifact changes
- teaching rules change
- the Editorial Brief changes
- a Reasoning Path is deprecated
- graph facts change
- references are updated
- diagnosis identity is merged
- generation policy changes
- the AI model or schema changes materially

Old drafts should not remain indefinitely applicable without renewed context validation.

44. Quality Standard for AI-Assisted Editorial Work

High-quality AI-assisted work should be:

Context-grounded

The draft visibly reflects governed diagnosis and educational context.

Purpose-specific

The output solves an identified editorial need.

Structurally valid

The draft follows the target artifact schema.

Clinically defensible

Claims and relationships are suitable for expert review.

Educationally useful

The draft supports a defined learner or editorial outcome.

Transparent

The source issue, context, output, and decisions are preserved.

Reviewable

The reviewer can compare the proposal with current governed content.

Reversible

The draft can be rejected, superseded, or removed without erasing history.

Controlled in application

Accepted content enters the system through an explicit versioned change.

Human-governed

No AI output independently receives institutional authority.

45. Normative Statements

- Every material AI contribution **MUST** preserve AI provenance.
- Every AI generation request **MUST** identify an editorial purpose.
- AI generation **MUST** use a governed diagnosis identity where available.
- Generation Context **SHOULD** include applicable learning goals, teaching rules, mimics, discriminators, pitfalls, and difficulty guidance.
- AI **MUST NOT** infer institutional approval from the presence of existing content.
- AI-generated content **MUST** begin as candidate or draft knowledge.
- AI-generated cases **MUST NOT** enter playable or published states directly.

- AI-generated Diagnosis Education **MUST NOT** enter approved or published states directly.
 - AI-generated Graph Candidates **MUST NOT** be promoted automatically into active Graph Facts.
 - AI-generated Teaching Relationships, Evidence Relationships, and Reasoning Paths **MUST NOT** become active without human review.
 - AI-generated differential mappings **MUST NOT** silently resolve clinically ambiguous diagnosis identities.
 - AI-generated clinical claims **MUST** remain distinguishable from verified evidence.
 - AI **MUST NOT** be represented as a clinical evidence source.
 - High-risk management content **MUST** receive heightened clinical review.
 - Validation scores and trust tiers **MUST NOT** be interpreted as approval.
 - A passed AI validation **MUST NOT** replace human editorial review.
 - Hallucination-risk signals **MUST** remain visible until reviewed or resolved.
 - Accepted AI Drafts **MUST NOT** be treated as applied automatically.
 - Application **MUST** produce an auditable artifact change.
 - Materially changed target artifacts **MUST** invalidate or trigger reassessment of pending drafts.
 - Applying an AI Draft **SHOULD** trigger appropriate revalidation.
 - Rejected and superseded AI Drafts **MUST** remain historically traceable.
 - AI Draft conflicts **MUST** be resolved explicitly.
 - Publication **MUST** apply to the governed artifact resulting from AI-assisted work, not to the AI Draft itself.
 - Human editors **MUST** retain final authority over approval and publication.
 - AI operational failure **MUST NOT** partially modify governed artifacts.
 - AI Drafts **SHOULD** expire or require reassessment when their generation context becomes stale.
-

Closing Principle

Artificial intelligence can expand Wardle's editorial capacity, but it cannot create institutional trust by itself.

Trust arises when AI contribution is bounded by governed context, expressed as a transparent candidate draft, tested through validation, evaluated by qualified human reviewers, and applied through a controlled and auditable revision.

Within WEOS, AI does not replace editorial judgement.

It makes more editorial work possible while leaving responsibility, authority, and accountability with people.